

Data Blocks: Hybrid OLTP and OLAP on Compressed Storage using both Vectorization and Compilation

Authors: Harald Lang, Tobias Mühlbauer, Florian Funke,
Peter Boncz, Thomas Neumann, Alfons Kemper
(June 2016, DOI: 10.1145/2882903.2882925, Conference: SIGMOD)

Contributed by:

Group 7
Meet Shah 202011047
Krunal Mehta 202011051

Guided by:

Prof. Minal Bhise
Mayank Patel
Kalgi Gandhi

1. Introduction

Online Transaction Processing (OLTP) databases are designed for online transaction queries like insert, update and delete. An example is banking sector or stock market-related queries where updating is frequent. Online Analytical Processing (OLAP) deals with Archival Data or Historical Data and a very small amount of transactions. The main thing of this experiment is that it talks about hybrid OLTP and OLAP system. Here data set is divided into two parts, hot, which is frequently accessed, and cold, which is infrequently accessed. Data Blocks are essentially compressed cold data chunks ^[1]. Further, to speed up Data Blocks scans, Positional Small Materialized Aggregates (PSMAs) - a lightweight index structure containing a concise lookup table along with a simple min/max attribute is used. The scan range can be decreased by the PSMA indexes even if the whole Data Block cannot be skipped during a scan.

1.1 Motivation

As data is increasing rapidly, data summarization has to be done to handle large amounts of data efficiently. Data summarization affects primary storage data reduction and will continue to play an important role in data reduction.

1.2 Objective

The main goal is to improve query performance in a hybrid OLAP-OLTP database system by data summarization without affecting query efficiency. For that, division of data into hot and cold data, compression of cold data, and PSMA Lookup Table implementation are the main objectives of the experiment.

1.3 Scope

The experiment is done on Relational Database considering both types of queries, OLTP (Online Transaction Processing) and OLAP (Online Analytical Processing).

2. Data Blocks Flow, Data Structures, and Algorithm

This section elaborates flow, data structures, and algorithm of data blocks technique.

2.1 Flow

Here firstly, we are loading the TPC-H dataset into our system. Then we are dividing Lineitem Table into hot and cold chunks and compressing cold chunks. Further, we are generating PSMA Lookup Table.

Later, after the loading of the queryset, if a query is required to access cold data, we will do Data Blocks scan to the range extracted from PSMA Lookup Table, uncompress and scan it. Otherwise, a scan will be performed on hot uncompressed data.

Lastly, the query will be executed and the output will be shown.

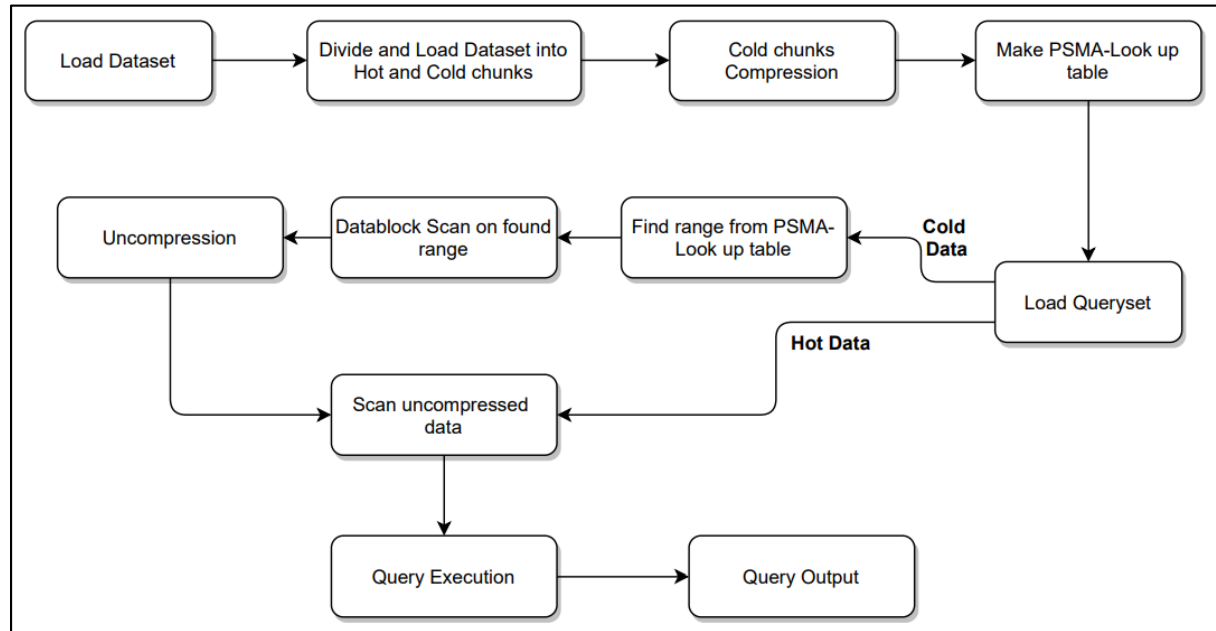


Fig.1. Flow of Data Blocks

2.2 Data Structures

In this experiment, a data-centric approach is used. Firstly Lineitem Table is sorted on the `l_shipdate` attribute, then 0.2M tuples were extracted from the starting date specified in the Key-Value List.

Further tuples are inserted into Lineitem_Hot Table and Lineitem_Cold Table according to the limit and offset values provided by Key-Value List.

For compression, truncation compression technique is applied on the `l_shipdate` attribute of Lineitem Cold Table.

PSMA Lookup Table will contain offset, limit, and no_of_records for each index.

Table1. Key-Value List

min_cold	0
max_cold	65536
min_hot	65536
max_hot	200000
partition_start_date	1992-01-02

Table2. Lineitem_Cold Table (Compressed)

l_orde rkey	l_part key	l_su ppke y	l_line numbe r	l_quan tity	l_exten dedpri ce	l_disc ount	l_ta x	l_retur nflag	l_line status
89173	4190	5	17	19756.89	0.04	0.01	A	F	89173
148071	3100	5	29	32453.03	0.05	0.06	A	F	148071
176498	1533	5	16	25191.84	0.06	0.02	R	F	176498
64313	1832	5	37	47260.47	0.06	0.02	R	F	64313
4641	4642	2	6	9273.84	0	0.08	A	F	4641

l_shipda te	l_commitdat e	l_receiptdate	l_shipinstr uct	l_shipmo de	l_comment
1	1994-12-26	1995-01-01	DELIVER IN PERSON	REG AIR	inal packages haggle carefully
11	1994-11-17	1995-02-04	COLLECT COD	MAIL	l requests hagg
12	1994-12-27	1995-01-12	DELIVER IN PERSON	SHIP	kages cajole carefully
14	1994-12-29	1995-01-31	DELIVER IN PERSON	FOB	deas use blithely! special foxes print af
19	1994-12-08	1995-02-10	NONE	TRUCK	sits wake furiously regular

Table3. Lineitem_Hot Table

l_orde rkey	l_part key	l_su ppke y	l_line numbe r	l_qua ntity	l_extend edprice	l_disc ount	l_ta x	l_retur nflag	l_line status
359	11158	6161	2	18	19244.7	0	0.03	A	F
391	121586	6611	1	14	22506.12	0.09	0.02	R	F
12101	1596	1597	2	25	37439.75	0.04	0.04	A	F
930	99635	2145	4	21	34327.23	0.06	0.02	A	F
10722	113068	5580	6	6	6486.36	0.09	0.03	R	F

l_shipdate	l_commitdat e	l_receiptdat e	l_shipinstruct	l_ship mode	l_comment
1995-01-27	1995-03-18	1995-01-31	DELIVER IN PERSON	RAIL	unusual warthogs. ironically sp
1995-02-11	1995-02-03	1995-02-13	TAKEBACK RETURN	TRUC K	escapades sleep furiously about
1995-02-13	1995-03-02	1995-02-14	COLLECT COD	REG AIR	al, final foxes about the ir
1995-02-16	1995-03-03	1995-03-13	DELIVER IN PERSON	SHIP	foxes. regular deposits integrate carefu
1995-02-12	1995-04-02	1995-02-21	NONE	FOB	eas. carefully special deposits after the

Table4. PSMA Lookup Table

index	offset	limit	no_of_records
0	0	2497	2497
1	2497	4979	2482
2	4979	7470	2491
3	7470	9910	2440
4	9910	12506	2596

2.3 Algorithm

Below algorithm is the step by step process in order to execute query by Data Blocks technique for hybrid OLAP-OLTP system.

1. Start
2. Load Dataset
3. Initialize variables according to Data Structures
4. Divide data into hot data and cold data
 - a. Sort Lineitem Table on l_shipdate attribute
 - b. Insert data into Lineitem_Cold Table and Lineitem_Hot Table according to Data Structures variables
5. Compress cold data using truncation compression scheme
 - a. Extract min and max of l_shipdate from Lineitem_Cold Table
 - b. Calculate and store delta value of l_shipdate for all tuples by subtracting min value from it
6. Make a PSMA Lookup Table for cold data
 - a. Find number of entry in PSMA Lookup Table by subtracting min l_shipdate from max l_shipdate
 - b. Insert index of each entry as delta value
 - c. Insert offset value of each entry by calculating number of records of all previous entries
 - d. Insert number_of_records value of each entry by calculating number of tuples (of that entry) present in Lineitem_Cold Table
 - e. Insert limit value of each entry by addition of offset value and number_of_records value
7. Load queryset
8. For all query in queryset do
 - If Data Block = cold
 - Find out range from PSMA Lookup Table
 - Data Block scan on found range
 - Uncompress data
 - Scan uncompressed data
 - Else
 - Scan uncompressed data
 - End if
 - Query Output
- End-for
9. Stop

3. Experimental Setup

This section contains details regarding the Dataset and Queryset used in the experiment, Software and Hardware specification for the execution of the experiment, and Evaluation Parameters of the experiment.

3.1 Dataset and Queryset

Benchmark TPC-H dataset is a decision support benchmark, it consists of a set of business oriented queries [2]. Size of the dataset used here is 21 MB which comprises 0.2M tuples.

3.2 Software and Hardware Setup

Experiments are performed using a system with 8 GB memory, core i5 processor, with a hard disk of 2TB. PostgreSQL 10.14 and Python 3.9 are used as software for the implementation [3, 4]. Additionally psycopg2 and matplotlib modules were used in Python in order to connect with PostgreSQL and to plot graphs respectively.

3.3 Evaluation Parameters

Evaluation Parameters for the experiment are identified as Algorithm Execution Time (AET) which calculates execution time of each algorithm. Execution time of each query is depicted as Query Execution Time (QET).

3.4 Block Diagram

Block Diagram of the experiment is shown in Fig.2.

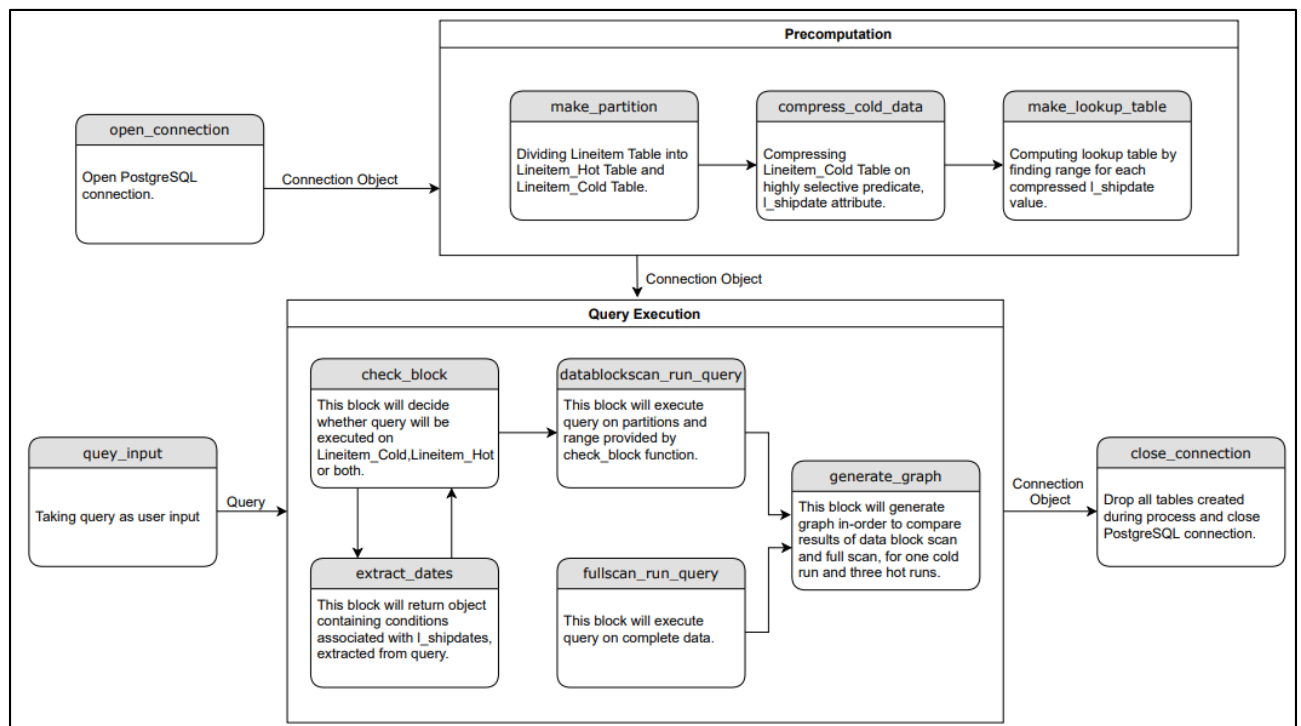


Fig.2. Block Diagram

4. Results and Discussions

This section contain the results obtain from the experiment.

4.1 PSMA Lookup Table Implementation

PSMA Lookup Table containing offset, limit, and no_of_records for each delta value as index is depicted in Fig.3.

C:\Windows\System32\cmd.exe - python latest.py

-> Creating Look up Table....
-> Total number of entries in look up table: 80

***** Look up table *****

Index	Offset	Limit	No_of_records
0	0	17	17
1	17	58	41
2	58	105	47
3	105	182	77
4	182	288	106
5	288	429	132
6	429	551	121
7	551	704	153
8	704	882	178
9	882	1084	202
10	1084	1334	250
11	1334	1611	277
12	1611	1898	287
13	1898	2183	285
14	2183	2494	311
15	2494	2845	351
16	2845	3195	350
17	3195	3560	365
18	3560	3921	361
19	3921	4312	391
20	4312	4723	411
21	4723	5148	425
22	5148	5634	486
23	5634	6121	487
24	6121	6663	542
25	6663	7228	565
26	7228	7798	570
27	7798	8325	527
28	8325	8894	569
29	8894	9524	630
30	9524	10189	665
31	10189	10841	652
32	10841	11512	671
33	11512	12211	699
34	12211	12946	735
35	12946	13681	735
36	13681	14417	736
37	14417	15199	782
38	15199	16001	802
39	16001	16846	845
40	16846	17699	853
41	17699	18518	819
42	18518	19388	869

Fig.3. PSMA Lookup Table

4.2 Query Execution

Query execution with Data Block Scan algorithm, and comparison with Full Scan is shown in Fig.4, Fig.5 and Fig.6.

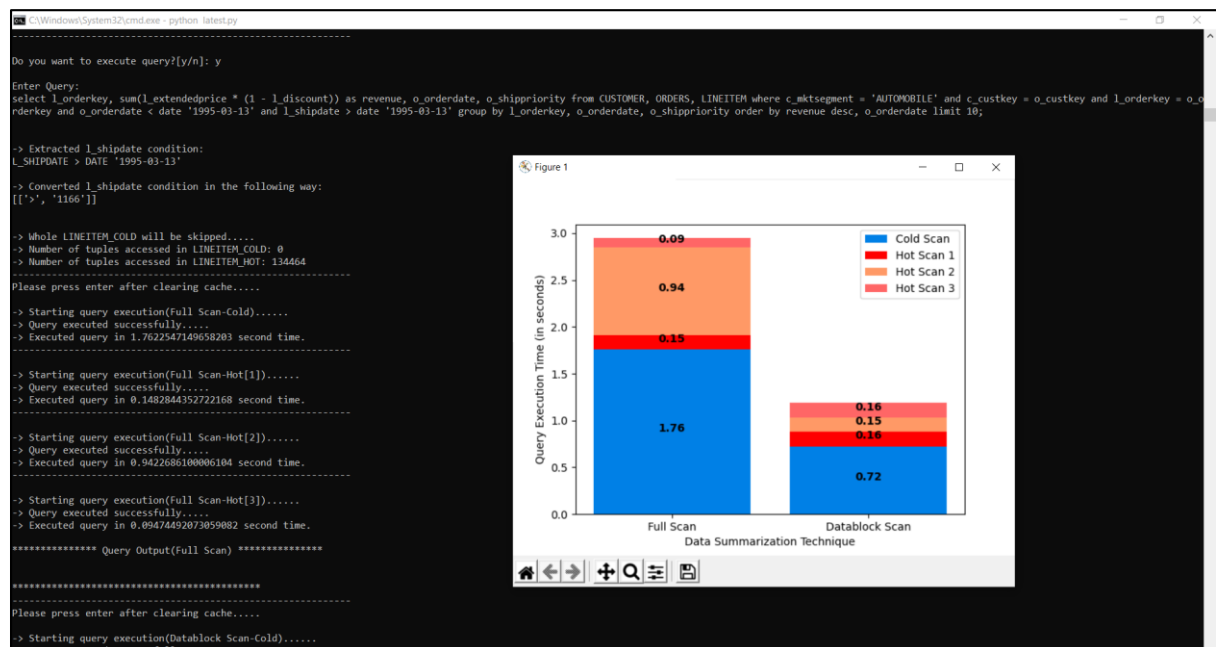


Fig.4. Query Execution (OLAP-Q3)

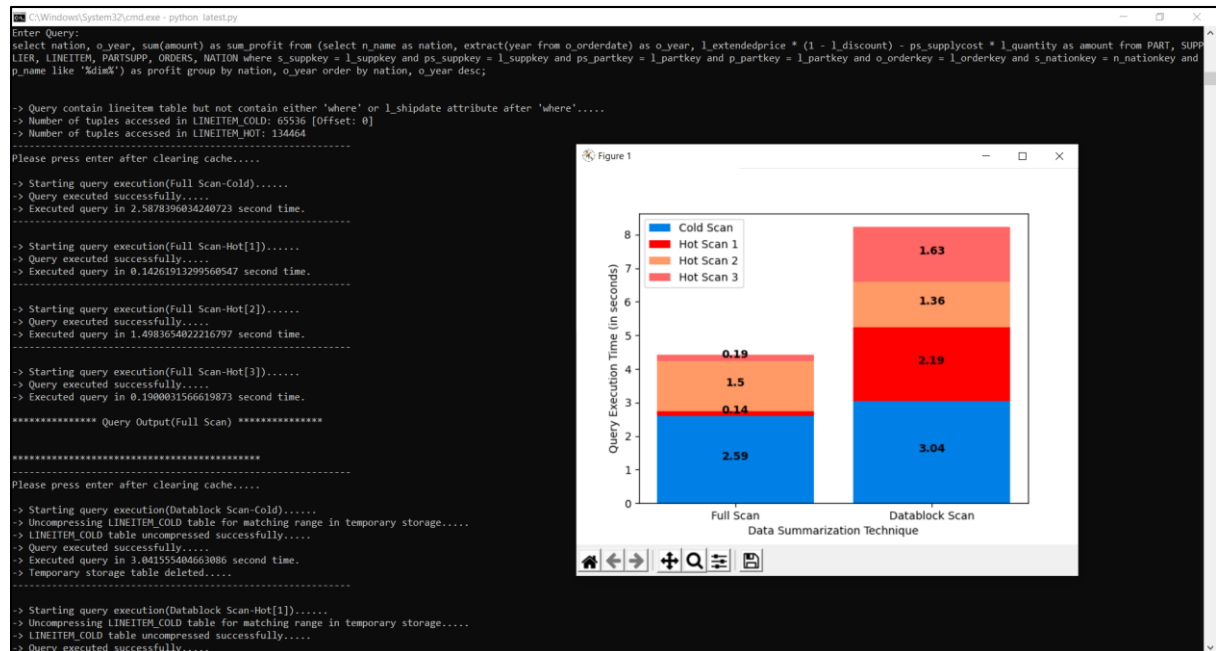


Fig.5. Query Execution (OLAP-Q9)

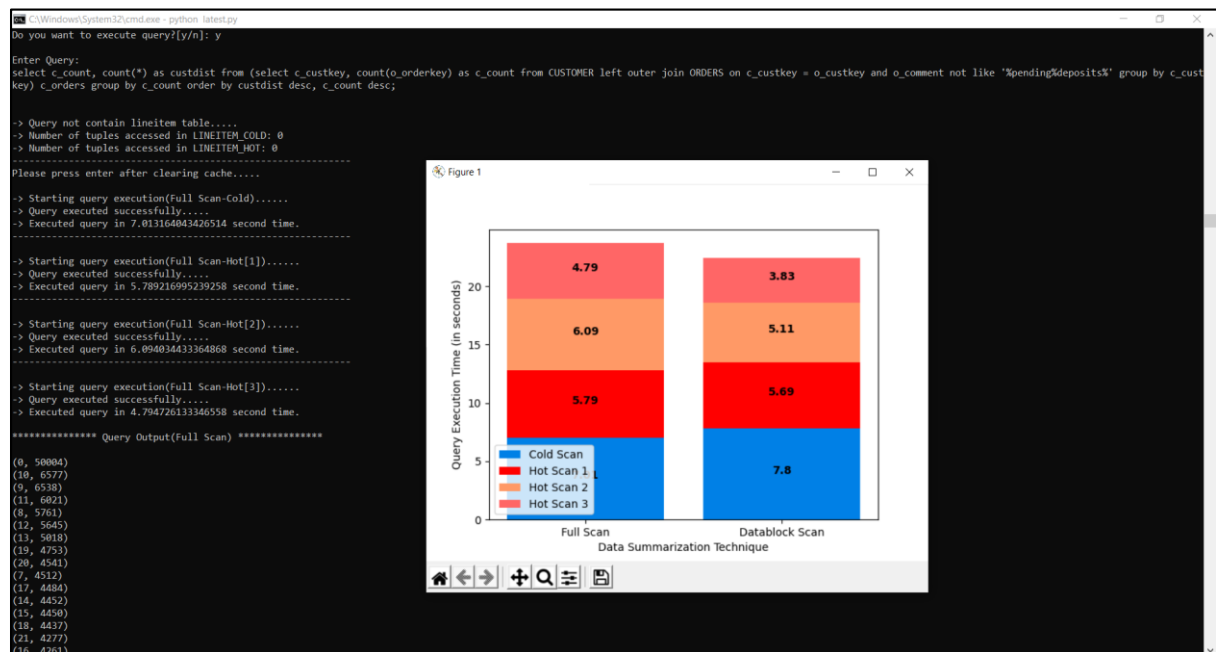


Fig.6. Query Execution (OLAP-Q13)

4.3 Experiment Results

AET of the Data Blocks experiment is recorded as 342.96 seconds.

For each query, algorithm is executed once for cold run and thrice for hot run, for both Full Scan and Data Block Scan.

4.3.1 Data Summarization (OLAP)

QET for OLAP is depicted in Fig.7, which shows average QET for both Full Scan and Data Block Scan for all 22 OLAP TPC-H benchmark queries.




Fig.7. Data Summarization (OLAP)

4.3.2 Data Summarization (OLTP)

QET for OLTP is depicted in Fig.8, which shows average QET for both Full Scan and Data Block Scan for all 22 modified OLTP TPC-H queries.




Fig.8. Data Summarization (OLTP)

5. Conclusion

We observe that when a query has to access Lineitem_Cold Table for evaluation, its overhead increase due to first uncompressing required cold data and then executing a query on that, but in the case when the query does not need to access Lineitem_Cold for evaluation it performs better as it skips Lineitem_Cold Table data and only executes a query on Lineitem_Hot Table data.

6. References

- [1] H. Lang, T. Muhlbauer, F. Funke, P. Boncz, T. Neumann, A. Kemper, Data Blocks: Hybrid OLTP and OLAP on Compressed Storage using both Vectorization and Compilation, SIGMOD, 2016, pp. 311-326.
- [2] TPC-H Dataset and Queryset: <http://www.tpc.org/tpch/>
- [3] Python: <https://www.python.org/downloads/windows/>
- [4] PostgreSQL: <https://www.postgresql.org/download/>